


University of Minho
School of Engineering

Romeu Alexandre Ribeiro Fernandes

Computational Discovery of DNA Aptamer-Graphene Sensors



University of Minho
School of Engineering

Romeu Alexandre Ribeiro Fernandes

Computational Discovery of DNA Aptamer-Graphene Sensors

Master's Dissertation in Bioinformatics

Dissertation supervised by

Dr. João Carneiro

Dr. Pedro Soares

Copyright and Terms of Use for Third Party Work

This dissertation reports on academic work that can be used by third parties as long as the internationally accepted standards and good practices are respected concerning copyright and related rights.

This work can thereafter be used under the terms established in the license below.

Readers needing authorization conditions not provided for in the indicated licensing should contact the author through the RepositóriUM of the University of Minho.

License granted to users of this work:

[Caso o autor pretenda usar uma das licenças Creative Commons, deve escolher e deixar apenas um dos seguintes ícones e respetivo lettering e URL, eliminando o texto em itálico que se lhe segue. Contudo, é possível optar por outro tipo de licença, devendo, nesse caso, ser incluída a informação necessária adaptando devidamente esta minuta]



CC BY

<https://creativecommons.org/licenses/by/4.0/> *[Esta licença permite que outros distribuam, remixem, adaptem e criem a partir do seu trabalho, mesmo para fins comerciais, desde que lhe atribuam o devido crédito pela criação original. É a licença mais flexível de todas as licenças disponíveis. É recomendada para maximizar a disseminação e uso dos materiais licenciados.]*



CC BY-SA

<https://creativecommons.org/licenses/by-sa/4.0/> *[Esta licença permite que outros remis-*

turem, adaptem e criem a partir do seu trabalho, mesmo para fins comerciais, desde que lhe atribuam o devido crédito e que licenciem as novas criações ao abrigo de termos idênticos. Esta licença costuma ser comparada com as licenças de software livre e de código aberto «copyleft». Todos os trabalhos novos baseados no seu terão a mesma licença, portanto quaisquer trabalhos derivados também permitirão o uso comercial. Esta é a licença usada pela Wikipédia e é recomendada para materiais que seriam beneficiados com a incorporação de conteúdos da Wikipédia e de outros projetos com licenciamento semelhante.]



CC BY-ND

<https://creativecommons.org/licenses/by-nd/4.0/> [Esta licença permite que outras pessoas usem o seu trabalho para qualquer fim, incluindo para fins comerciais. Contudo, o trabalho, na forma adaptada, não poderá ser partilhado com outras pessoas e têm que lhe ser atribuídos os devidos créditos.]



CC BY-NC

<https://creativecommons.org/licenses/by-nc/4.0/> [Esta licença permite que outros remisturem, adaptem e criem a partir do seu trabalho para fins não comerciais, e embora os novos trabalhos tenham de lhe atribuir o devido crédito e não possam ser usados para fins comerciais, eles não têm de licenciar esses trabalhos derivados ao abrigo dos mesmos termos.]



CC BY-NC-SA

<https://creativecommons.org/licenses/by-nc-sa/4.0/> [Esta licença permite que outros remisturem, adaptem e criem a partir do seu trabalho para fins não comerciais, desde que lhe atribuam a si o devido crédito e que licenciem as novas criações ao abrigo de termos idênticos.]



CC BY-NC-ND

<https://creativecommons.org/licenses/by-nc-nd/4.0/> [Esta é a mais restritiva das nossas seis licenças principais, só permitindo que outros façam download dos seus trabalhos e os partilhem desde que lhe sejam atribuídos a si os devidos créditos, mas sem que possam alterá-los de nenhuma forma ou utilizá-los para fins comerciais.]

Acknowledgements

Statement of Integrity

I hereby declare having conducted this academic work with integrity.

I confirm that I have not used plagiarism or any form of undue use of information or falsification of results along the process leading to its elaboration.

I further declare that I have fully acknowledged the Code of Ethical Conduct of the University of Minho.

University of Minho, Braga, august 2026

Romeu Alexandre Ribeiro Fernandes

Abstract

Rapid and reliable detection of microbial DNA is central to clinical diagnostics, antimicrobial-resistance surveillance and environmental monitoring, yet the dominant amplification-based methods require laboratory infrastructure that limits their use at the point of care. Graphene field-effect transistor (GFET) biosensors, functionalised with single-stranded DNA recognition elements, offer a compelling label-free alternative, converting the hybridisation of a target sequence at the graphene surface into a measurable electrical signal. The bottleneck of this approach is the recognition sequence: identifying, for each pathogen, probe candidates that are conserved across strains, hybridise favourably, and are free of stable self-structure, without resorting to costly and low-throughput experimental screening of hundreds of candidates.

This dissertation addresses that bottleneck with a reproducible computational pipeline for the *discovery* of DNA probe candidates for GFET biosensors, developed for a panel of clinically relevant bacterial marker genes. From sequences retrieved from the NCBI databases, the pipeline performs automatic length clustering, multiple sequence alignment with MAFFT, and per-column conservation scoring based on the Percentage of Pairwise Identity. Sliding-window candidates are screened with nearest-neighbour thermodynamics (melting temperature, GC content, hairpin and homodimer free energies) and evaluated for secondary structure, condensed into a continuous “No-fold” score. All thresholds are resolved through an organism-aware, layered scheme (gene, species and organism type), so that pathogens with very different genome composition are judged fairly. The highest-quality candidates are prioritised and validated three-dimensionally, and a curated, feature-rich dataset of the best candidates per gene is produced.

Applied to the six-gene panel, the pipeline reduced 8455 candidate windows to 2943 that passed thermodynamic and structural screening, from which the best probes per gene were validated in 3D. Alignment-free diversity and rarefaction analyses characterised each target and justified the sampling depth. The result is a prioritised, biophysically sound set of probe candidates and a reproducible framework that provides the sequence-level foundation for the rational design of aptamer–graphene GFET biosensors, and that generalises to new genes and organisms with minimal configuration.

Keywords DNA aptamer, DNA probe, graphene biosensor, graphene field-effect transistor (GFET), label-free detection, microbial DNA detection, point-of-care diagnostics, probe design, sequence conservation, multiple sequence alignment, melting temperature, secondary structure, No-fold score, 3D structure prediction

Resumo

A detecção rápida e fiável de DNA microbiano é central para o diagnóstico clínico, a vigilância da resistência antimicrobiana e a monitorização ambiental, mas os métodos dominantes, baseados em amplificação, exigem infraestrutura laboratorial que limita o seu uso no local de prestação de cuidados. Os biossensores baseados em transístores de efeito de campo de grafeno (GFET), funcionalizados com elementos de reconhecimento de DNA de cadeia simples, oferecem uma alternativa sem marcação apelativa, convertendo a hibridação de uma sequência-alvo à superfície do grafeno num sinal elétrico mensurável. O estrangulamento desta abordagem é a sequência de reconhecimento: identificar, para cada agente patogénico, sondas candidatas que sejam conservadas entre estirpes, hibridem favoravelmente e estejam livres de estrutura secundária estável, sem recorrer à triagem experimental dispendiosa e de baixo débito de centenas de candidatas.

Esta dissertação aborda esse estrangulamento com um pipeline computacional reproduzível para a *descoberta* de sondas de DNA candidatas para biossensores GFET, desenvolvido para um painel de genes marcadores bacterianos clinicamente relevantes. A partir de sequências obtidas nas bases de dados do NCBI, o pipeline realiza a seleção automática do agrupamento de comprimento, o alinhamento múltiplo com o MAFFT e a pontuação de conservação por coluna baseada na Percentagem de Identidade Par-a-Par. As candidatas, geradas por janela deslizante, são filtradas com termodinâmica de vizinhos mais próximos (temperatura de fusão, conteúdo em GC, energias de hairpin e homodímero) e avaliadas quanto à estrutura secundária, condensada num score contínuo “No-fold”. Todos os limiares são resolvidos através de um esquema em camadas e sensível ao organismo (gene, espécie e tipo de organismo), para que agentes patogénicos com composição genómica muito distinta sejam avaliados de forma justa. As candidatas de maior qualidade são priorizadas e validadas tridimensionalmente, e é produzido um conjunto de dados curado e rico em atributos com as melhores candidatas por gene.

Aplicado ao painel de seis genes, o pipeline reduziu 8455 janelas candidatas a 2943 que passaram a triagem termodinâmica e estrutural, a partir das quais as melhores sondas por gene foram validadas em 3D. Análises de diversidade sem alinhamento e de rarefação caracterizaram cada alvo e justificaram

a profundidade de amostragem. O resultado é um conjunto priorizado e biofisicamente sólido de sondas candidatas e uma estrutura reprodutível que constitui o alicerce, ao nível da sequência, para o desenho racional de biossensores GFET de aptâmero-grafeno, generalizável a novos genes e organismos com configuração mínima.

Palavras-chave aptâmero de DNA, sonda de DNA, biossensor de grafeno, transistor de efeito de campo de grafeno (GFET), detecção sem marcação, detecção de DNA microbiano, diagnóstico no local de prestação de cuidados, desenho de sondas, conservação de sequência, alinhamento múltiplo de sequências, temperatura de fusão, estrutura secundária, score No-fold, previsão de estrutura 3D

Contents

1	Introduction	1
1.1	Context	1
1.2	Problem statement	2
1.3	Objectives	2
1.4	Target panel	3
1.5	Structure of the dissertation	4
2	State of the Art	5
2.1	DNA aptamers and probes as recognition elements	5
2.2	Graphene GFET biosensors for DNA detection	5
2.3	The DNA–graphene interface	6
2.4	Computational design and validation of probe sequences	6
2.5	Summary	7
3	Methods	8
3.1	Sequence retrieval from NCBI	8
3.2	Length clustering	8
3.3	Multiple sequence alignment	9
3.4	Conservation scoring (PPI)	9
3.5	Candidate generation	9
3.6	Thermodynamic screening	10
3.7	Secondary structure and the No-fold score	10
3.8	Organism-aware parameters	11
3.9	Quality-based selection	12
3.10	Three-dimensional validation	12
3.11	Assembly of probe–graphene models	13

3.11.1	Generation of the graphene sheet with GOPY	13
3.11.2	Assembly of the probe–graphene models	14
3.12	Integration of a reference probe panel	15
3.13	Diversity and rarefaction analysis	15
3.14	Implementation and reproducibility	16
4	Results and Discussion	17
4.1	Retrieval and length curation	17
4.2	Screening funnel	18
4.3	Three-dimensional validation	18
4.4	Genetic diversity of the targets	19
4.5	Choice of sampling depth	20
4.6	Effect of the organism-aware parameters	20
4.7	Comparison with the reference panel	21
4.8	Probe–graphene models	21
4.9	Curated dataset of prioritised candidates	24
5	Conclusions and Future Work	25
5.1	Conclusions	25
5.2	Prospects for future work	26

List of Figures

1	Distribution of the <code>seqfold</code> secondary-structure free energy (ΔG_{MFE}) over the candidate probes, from which the -2.6 kcal/mol threshold (fifth percentile) is taken.	11
2	Pairwise 4-mer distance heatmaps per gene (alignment-free). Each cell is the 4-mer distance between two retrieved sequences of the gene.	20
3	Rarefaction of the mean 4-mer diversity (left) and of the richness, i.e. the number of distinct 4-mers (right), per gene, as a function of the number of subsampled sequences.	20
4	Assembled probe–graphene models for <i>p5972</i> (<i>algD</i> , top) and <i>p7716</i> (<i>frdB</i> , bottom), rendered from the model coordinates. Left: side view, with the graphene plane, the probe above it and the $3.2 \text{ \AA}$ separation indicated. Right: top view, showing the probe footprint within the $121.6 \times 119.8 \text{ \AA}$ sheet. Probe atoms are coloured by element; graphene carbons are grey.	23

List of Tables

1	Target genes and their host organisms used as the working panel.	4
2	Global default thresholds. Values are overridden per organism type, species or gene where justified (Section 3.8).	10
3	Per-type parameter profiles. Viruses use a wider GC window and a more permissive conservation threshold; the protozoan profile is AT-shifted.	12
4	Sequences retrieved and retained per gene, and the resulting alignment width.	17
5	Screening funnel: candidate windows surviving each successive filter. “Basic” is the number passing all four thermodynamic criteria; “Seqfold” additionally passes the secondary-structure screen.	18
6	Top probe candidates, strong in both sequence quality and 3D confidence.	19
7	Alignment-free 4-mer diversity per gene. Bold marks the most diverse (<i>lytA</i>) and the most conserved/redundant (<i>frdB</i>) genes.	19
8	Geometry of the 30 assembled probe–graphene models. “Extent” is the largest in-plane dimension of the probe (which determined the sheet size), “height” its vertical extent above the sheet, and “closest contact” the shortest probe–sheet interatomic distance. . .	22

Acronyms

T_m melting temperature.

AMR Antimicrobial Resistance.

FRET Förster Resonance Energy Transfer.

GC Guanine-Cytosine.

GFET Graphene Field-Effect Transistor.

GO Graphene Oxide.

GPU Graphics Processing Unit.

IUPAC International Union of Pure and Applied Chemistry.

MFE Minimum Free Energy.

MSA Multiple Sequence Alignment.

NCBI National Center for Biotechnology Information.

NN Nearest-Neighbour.

PAE Predicted Aligned Error.

PCR Polymerase Chain Reaction.

pLDDT predicted Local Distance Difference Test.

PoC Point-of-Care.

PPI Percentage of Pairwise Identity.

SELEX Systematic Evolution of Ligands by EXponential enrichment.

ssDNA single-stranded DNA.

Chapter 1

Introduction

1.1 Context

Bacterial infections are among the leading causes of death worldwide, and the spread of **Antimicrobial Resistance (AMR)** is steadily eroding the ability to treat them: an estimated 4.95 million deaths were associated with bacterial **AMR** in 2019, of which 1.27 million were directly attributable to it (Murray et al., 2022). Rapid identification of the infecting pathogen, and increasingly of its resistance determinants, is therefore central to guiding therapy, to infection control and to public-health surveillance; several of the species targeted in this work feature on the WHO Bacterial Priority Pathogens List (World Health Organization, 2024).

The fast identification of microbial DNA underpins the diagnosis of infection and the tracking of environmental contamination. Culture-based methods are slow, and although the **Polymerase Chain Reaction (PCR)** and related amplification techniques are highly sensitive and specific, they require dedicated equipment, controlled laboratory conditions and lengthy sample preparation, which restricts their use outside the laboratory (Hu et al., 2024). This has driven the development of miniaturised solid-state biosensors, which convert a molecular-recognition event directly into an electrical or electrochemical signal and thereby enable rapid, portable and **Point-of-Care (PoC)** testing (Kozel and Burnham-Marusich, 2017; Krishnan et al., 2023).

Coupling DNA recognition elements to a **Graphene Field-Effect Transistor (GFET)** has been proposed as a particularly promising route to label-free nucleic-acid detection (Green and Norton, 2015). Graphene, a two-dimensional sheet of carbon isolated in 2004 (Novoselov et al., 2004), has its entire conduction channel exposed at its surface and transduces local charge variations with such sensitivity that even a single adsorbed molecule can be detected (Schedin et al., 2007); when it is functionalised with a **single-stranded DNA (ssDNA)** probe, the hybridisation of a complementary microbial target modulates the transistor current in a label-free manner. In practice, the probe is anchored to the graphene

surface through a linker, such as a pyrene derivative that binds non-covalently by π - π stacking, and target hybridisation shifts the transistor's charge-neutrality (Dirac) point; electrolyte-gated graphene devices of this kind have achieved label-free detection of DNA hybridisation down to the attomolar range (Campos et al., 2019). The performance of such devices, and the influence of the sensing interface and ionic conditions on it, is reviewed in detail in Chapter 2.

The specificity and sensitivity of such a device ultimately depend on the immobilised recognition sequence, which must meet several partly competing requirements: it must target a genomic region that is conserved across the clinical strains of the pathogen, so that strain variation does not cause false negatives; it must hybridise strongly and specifically at the assay temperature, which constrains its length, melting temperature and GC content; and it must avoid folding into stable hairpins or self-dimers that would compete with hybridisation. These requirements are biophysical and quantifiable, which makes them amenable to systematic computational evaluation.

Building on these experimental developments, a central practical obstacle remains: the high cost and low throughput of experimentally testing hundreds of candidate probes in order to discriminate multiple pathogenic species. This motivates a *complementary computational screening strategy* that can prioritise a small number of promising candidates before any experimental effort is spent.

1.2 Problem statement

The scientific problem addressed in this dissertation is how to computationally *discover and prioritise* DNA probe candidates that recognise specific microbial DNA sequences and are suitable for transduction on a GFET. This requires analysing the sequence and structural features — genomic conservation, hybridisation thermodynamics and self-structure — that determine whether a probe will bind its target strongly and specifically across clinical strains, and doing so in a reproducible, automated way that scales across many genes and organisms of very different genome composition (Green and Norton, 2015; Campos et al., 2019).

1.3 Objectives

The main objective is to build a reproducible computational pipeline that turns a target gene into a small set of ranked, biophysically validated ssDNA probe candidates for GFET biosensing. This decomposes into the following specific objectives:

- select a panel of clinically relevant microbial targets and marker genes, and automate the retrieval and curation of their sequences from the [NCBI](#) databases;
- identify conserved regions through multiple sequence alignment and per-column conservation scoring, so that the probes detect their target across clinical strains;
- screen candidate probes with well-founded biophysical criteria (melting temperature, GC content, hairpin, homodimer and secondary-structure stability) and combine them into transparent quality scores;
- make the design criteria *organism-aware*, so that targets with very different genome composition (AT-rich or GC-rich) are evaluated fairly;
- prioritise and validate the most promising candidates by three-dimensional structure prediction; and
- produce curated, feature-rich datasets of the prioritised probe candidates.

1.4 Target panel

Six marker genes from six clinically important bacterial species were used as the working panel throughout this work (Table 1). Several are established targets for the molecular detection of their species: the thermonuclease gene *nuc* for *Staphylococcus aureus* ([Brakstad et al., 1992](#)), the autolysin gene *lytA* for *Streptococcus pneumoniae* ([Carvalho et al., 2007](#)), and the outer-membrane lipoprotein gene *oprL* for *Pseudomonas aeruginosa* ([De Vos et al., 1997](#)). The panel spans a wide range of genome composition — from AT-rich (*Staphylococcus aureus*, *Haemophilus influenzae*) to GC-rich (*Pseudomonas aeruginosa*) — which makes it a demanding test of a general-purpose design method.

Gene	Organism	Group	Function
<i>nuc</i>	<i>Staphylococcus aureus</i>	A	thermonuclease
<i>rmpM</i>	<i>Neisseria meningitidis</i>	A	class-4 outer-membrane protein
<i>lytA</i>	<i>Streptococcus pneumoniae</i>	B	autolysin
<i>oprL</i>	<i>Pseudomonas aeruginosa</i>	B	peptidoglycan-associated lipoprotein
<i>algD</i>	<i>Pseudomonas aeruginosa</i>	B	GDP-mannose 6-dehydrogenase
<i>frdB</i>	<i>Haemophilus influenzae</i>	B	fumarate reductase Fe-S subunit

Table 1: Target genes and their host organisms used as the working panel.

1.5 Structure of the dissertation

The remainder of this dissertation is organised as follows. Chapter 2 reviews the state of the art in DNA aptamers and probes, in graphene **GFET** biosensors for nucleic-acid detection, and in the computational methods used to design and validate probe sequences. Chapter 3 describes the methods, detailing every stage of the pipeline and the rationale for each design decision and threshold. Chapter 4 presents and discusses the results for the target panel — the screening funnel, the three-dimensional validation, the diversity and rarefaction analyses, and the comparison with a reference probe panel. Chapter 5 draws the main conclusions and outlines future work.

Chapter 2

State of the Art

This chapter reviews the background needed to frame the work. It covers DNA aptamers and probes as recognition elements, graphene **GFET** biosensors for nucleic-acid detection, the physics of the DNA–graphene interface, and the computational methods used to design and validate probe sequences.

2.1 DNA aptamers and probes as recognition elements

DNA aptamers are short single-stranded nucleic acids, selected *in vitro* by **Systematic Evolution of Ligands by EXponential enrichment (SELEX)** (Tuerk and Gold, 1990; Ellington and Szostak, 1990), that fold into defined structures capable of high-affinity, high-specificity recognition of a target (Stephens et al., 2025; Sun et al., 2022). Cell-**SELEX** and related variants allow aptamers to be raised against whole microbial cells and complex targets, yielding recognition elements for bacterial surface antigens, toxins and viral components (Kolm et al., 2020; Hu et al., 2024). Because they are chemically synthesised, aptamers and DNA probes are inexpensive, reproducible and easily modified, which makes them well suited to solid-state biosensors and to *in silico* optimisation (Sun et al., 2022). For the detection of a known microbial DNA sequence, a complementary hybridisation probe is the natural recognition element: it targets a defined genomic region and is captured by Watson–Crick pairing, which is the mechanism exploited throughout this work.

2.2 Graphene GFET biosensors for DNA detection

Graphene combines very high carrier mobility with a fully surface-exposed conduction channel, so that charges binding at its surface strongly modulate its conductance (Novoselov et al., 2004); the adsorption of even a single molecule can be resolved electrically (Schedin et al., 2007), making it an outstanding transducer for label-free sensing (Krishnan et al., 2023; Green and Norton, 2015), including for nucleic

acids at very low concentrations (Hwang et al., 2020). In a **Graphene Field-Effect Transistor (GFET)**, the graphene channel is gated electrolytically, and the hybridisation of a target strand to an immobilised probe shifts the transfer characteristic through the added backbone charge. Electrolyte-gated **GFETs** functionalised with DNA probes have reached attomolar, label-free detection of hybridisation (Campos et al., 2019), back-gated architectures have enabled scalable device production (Ping et al., 2016), and microfluidic integration has turned such devices into programmable lab-on-chip DNA sensors (Purwidyantri et al., 2022). A key operating constraint is the ionic strength of the medium: high salt concentration produces strong Debye screening that shortens the effective sensing length and attenuates the signal from charges far from the surface (Purwidyantri et al., 2021); transduction schemes based on target-induced conformational change of the probe have been shown to mitigate this limitation (Nakatsuka et al., 2018). These experimental studies define the device context that the computational design in this dissertation aims to support.

2.3 The DNA–graphene interface

The behaviour of a **GFET** DNA sensor is governed by how DNA interacts with graphene. Single-stranded DNA adsorbs strongly and flat onto graphene through π – π stacking of its exposed nucleobases, whereas double-stranded DNA binds much more weakly because base pairing sequesters the bases (Green and Norton, 2015; Krishnan et al., 2023). This contrast is the basis of many transduction schemes: for example, graphene-oxide **Förster Resonance Energy Transfer (FRET)** aptasensors, in which a single-stranded aptamer adsorbed on **Graphene Oxide (GO)** is desorbed upon target binding, restoring fluorescence (Lu et al., 2009; Zhu et al., 2015). Adsorption can, however, perturb DNA folding and kinetics, which matters especially for structured recognition elements: a useful probe must retain its ability to hybridise its target while still producing a detectable electronic perturbation at the surface (Green and Norton, 2015).

2.4 Computational design and validation of probe sequences

Because experimentally screening large numbers of candidate probes is costly and slow, computational methods that link sequence and structure to binding and transduction are increasingly used to design and prioritise candidates (Sun et al., 2022). Several families of criteria are well established. The stability of the probe–target duplex is quantified by its melting temperature (T_m), predicted accurately with the nearest-neighbour thermodynamic model (SantaLucia Jr. and Hicks, 2004), together with GC content and probe length (Wetmur, 1991); tools such as primer3 implement these calculations (Untergasser

et al., 2012). Self-structure — hairpins, self-dimers and the overall **Minimum Free Energy (MFE)** of the single strand — must be minimised, since a folded probe is unavailable for capture; these energies are obtained from nearest-neighbour folding models (Turner and Mathews, 2010; Zadeh et al., 2011; Timmons, 2020). To detect a pathogen across its strains, conserved genomic regions are identified from a **Multiple Sequence Alignment (MSA)** of many sequences, for which MAFFT is a fast and accurate aligner (Katoh and Standley, 2013); conservation is then scored per column, for example by the **Percentage of Pairwise Identity (PPI)** adopted from the *VirusScope* tool (Lima, 2023).

Beyond sequence, deep-learning methods now predict three-dimensional structure with high accuracy. AlphaFold achieved near-experimental protein predictions (Jumper et al., 2021), and AlphaFold 3 extended this to nucleic-acid-containing complexes with per-residue and pairwise confidence estimates (Abramson et al., 2024); open models such as Boltz-1/2 make these predictions freely runnable on commodity GPUs (Wohlwend et al., 2024). Although calibrated primarily for proteins, the predicted structure and the associated confidence of a short **ssDNA** probe provide a useful, independent check on a computationally selected sequence.

2.5 Summary

Graphene **GFET** devices are an excellent, experimentally proven transducer for label-free DNA detection, but their performance hinges on the recognition sequence. The biophysical criteria for a good probe — conservation, favourable hybridisation thermodynamics and absence of stable self-structure — are individually well understood and supported by mature software, yet they are rarely combined into a single, reproducible and organism-aware pipeline that also ranks the candidates and validates them structurally. The methodology presented in the next chapter is designed to address that gap.

Chapter 3

Methods

The probe-discovery method was implemented as a single, reproducible Python pipeline. Given a target gene, the pipeline proceeds through the stages described below. It is built on Biopython ([NCBI](#) access and sequence handling), an external MAFFT executable (alignment), `primer3-py` (thermodynamics), `seqfold` (secondary structure) and `pandas` (tabular output); the three-dimensional validation is run separately with Boltz-2 on a [Graphics Processing Unit \(GPU\)](#). Every design threshold is resolved through a layered scheme that can be overridden per gene, per species and per organism type (Section [3.8](#)).

3.1 Sequence retrieval from NCBI

For each target, sequences are retrieved from the [NCBI](#) nucleotide database through the Entrez interface. An `esearch` query returns up to a configurable number of identifiers (set to 100 in this work; justified in Section [3.13](#)), which are then fetched with `efetch` in FASTA format, in batches of 200 identifiers, with short pauses between requests to respect the [NCBI](#) rate limit of three requests per second without an API key. Each target defines a primary query combining the gene name, the organism and a sequence-length range; if it returns fewer than two records, the pipeline falls back in turn to one or two broader alternative queries and, failing those, to a fixed RefSeq accession. Records outside a coarse per-gene length range are discarded at this stage. For genes with a single usable record (allowed only for *mpm*), a single-sequence mode is used.

3.2 Length clustering

Public records mix full-length genes with partial fragments and oversized entries, which together produce gap-ridden alignments. An automatic length-clustering step therefore retains only the dominant length cluster: the sequence length whose $\pm 25\%$ band contains the most records is taken as the centre, and

only records within that band are kept. The selection is data-driven — no per-gene minimum or maximum length is tuned by hand — so the method transfers to new genes without adjustment. It is applied only when at least four records are available, and a safeguard prevents the set from being reduced below two sequences.

3.3 Multiple sequence alignment

The retained sequences are aligned with MAFFT in its automatic mode (`--auto`). When only one sequence is available the alignment step is bypassed and that sequence is used directly.

3.4 Conservation scoring (PPI)

Conservation is scored per alignment column using the **Percentage of Pairwise Identity (PPI)**, adopted from the *VirusScope* tool (Lima, 2023). For a column, the **PPI** is the average identity over all pairs of non-gap bases, computed from the base counts; pairs of ambiguous **IUPAC** bases contribute a partial weight (for example a purine code *R* scores 0.5 against *A* or *G*, and *N* scores 0.25 against any base). Columns in which fewer than 30% of the sequences have a non-gap base are assigned zero, and only columns with at least two non-gap bases are counted. Because the per-column **PPI** is independent of the window, it is pre-computed once per alignment, which is essential for hundreds of sequences. The conservation of a candidate window is the mean of the **PPI** of its valid columns. In contrast to primer design, the identity of the terminal 3' nucleotides (scored separately as PPI3 in the original tool) is deliberately excluded, because capture probes are not extended by a polymerase and so the fidelity of their 3' end is irrelevant.

3.5 Candidate generation

Candidate probes are generated by sliding a window over the alignment. For each start position, taken every 3 columns, windows of every length from 18 to 28 nucleotides are formed. A window is discarded if any of its columns has a gap fraction above the allowed maximum (20% by default), or if its mean conservation is below the target's conservation threshold. The sequence of a surviving window is the per-column consensus (the most frequent non-gap base).

3.6 Thermodynamic screening

Each candidate is scored with `primer3-py` under fixed conditions: a monovalent cation concentration of 50 mM, no divalent cations or dNTPs, an oligonucleotide concentration of 250 nM and a reference temperature of 37 °C. Four quantities are computed: the melting temperature T_m (**Nearest-Neighbour (NN)** model (SantaLucia Jr. and Hicks, 2004)), the GC content (the fraction of G and C bases), the hairpin free energy and the homodimer (self-dimer) free energy, the last two reported in kcal/mol and only when a structure is found. A candidate passes the basic screen when its T_m and GC content fall within the resolved windows and its hairpin and homodimer free energies are no more negative than their respective limits (Table 2). The T_m window may optionally be derived from the hybridisation assay temperature T as $[T + 15, T + 35]$ °C.

Parameter	Default value	Basis
Probe length	18–28 nt	Wetmur (1991)
Melting temperature T_m	53–72 °C	SantaLucia Jr. and Hicks (2004); assay-derivable
GC content	40–60%	IDT OligoAnalyzer (Integrated DNA Technologies)
Hairpin ΔG	≥ -2.0 kcal/mol	Integrated DNA Technologies; primer3 (Untergasser et al., 2012)
Homodimer ΔG	≥ -5.0 kcal/mol	IDT OligoAnalyzer (Integrated DNA Technologies)
Conservation (PPI)	≥ 0.85	pan-strain detection (Lima, 2023)
Alignment gap fraction	$\leq 20\%$	alignment quality
Secondary structure ΔG_{MFE}	≥ -2.6 kcal/mol	5th percentile of the data

Table 2: Global default thresholds. Values are overridden per organism type, species or gene where justified (Section 3.8).

3.7 Secondary structure and the No-fold score

The single-strand secondary structure of each candidate that passed the basic screen is computed with `seqfold`, which returns the **Minimum Free Energy (MFE)** at 37 °C and the fraction of paired bases. A candidate passes the secondary-structure screen when its **MFE** is no more negative than -2.6 kcal/mol; this value is the fifth percentile of the observed **MFE** distribution (Figure 1), and fixing it keeps the criterion reproducible across runs rather than always cutting a fixed fraction of the data.




Figure 1: Distribution of the seqfold secondary-structure free energy (ΔG_{MFE}) over the candidate probes, from which the -2.6 kcal/mol threshold (fifth percentile) is taken.

In addition to this pass/fail screen, the three folding energies (hairpin, homodimer and **MFE**) are combined into a continuous *No-fold* score in the range 0–100, following *VirusScope* (Lima, 2023). Each energy is mapped through a logistic penalty,

$$\text{No-fold} = \left(1 - \frac{1}{1 + \exp(x)} \right) \times 100, \quad x = \frac{\Delta G_{\text{cal}} - \Delta G_0/2}{s R T},$$

with $\Delta G_0 = -7000$ cal/mol, slope $s = 1$, $R = 1.987$ cal mol $^{-1}$ K $^{-1}$ and $T = 310$ K; the per-probe No-fold score is the mean of the scores of the energies that are available. A high No-fold score denotes a probe that is largely free of stable self-structure. This continuous score is used, together with conservation, to rank the candidates rather than merely to filter them.

3.8 Organism-aware parameters

Because the intended application spans pathogens with very different genomes, a single global set of thresholds is inadequate: criteria appropriate for a balanced bacterial genome would unfairly reject an AT-rich or a GC-rich target. All thresholds are therefore resolved from the most specific layer to the most general,

$$\text{gene} \rightarrow \text{species} \rightarrow \text{organism type} \rightarrow \text{global default},$$

where the first layer that defines a parameter wins. Five organism types are distinguished, differing mainly in genome GC content and in variability (Table 3); species-level profiles then fine-tune the values by the GC content of each genome (for example *P. aeruginosa* at $\sim 67\%$ GC (Stover et al., 2000), *A. fumigatus* at $\sim 50\%$ GC (Nierman et al., 2005) and the AT-rich *P. falciparum* at $\sim 19\%$ GC (Gardner et al., 2002)).

Organism type	GC window	Conservation	Max length
Bacteria	40–60%	0.85	28
Virus	35–68%	0.70	30
Fungus	40–60%	0.80	28
Protozoan	15–40%	0.80	30
Host (human)	35–60%	0.85	28

Table 3: Per-type parameter profiles. Viruses use a wider GC window and a more permissive conservation threshold; the protozoan profile is AT-shifted.

When the pipeline encounters an organism for which no profile exists and is run interactively, it prompts the user for the organism type and for each threshold, proposing a suggested value that can be accepted or overridden; the resulting profile is stored and reused in later runs. When run non-interactively, it applies the suggested values without prompting.

3.9 Quality-based selection

Because three-dimensional prediction is comparatively expensive, only the best candidates per gene are carried forward. Rather than ranking by melting temperature — which the thermodynamic screen has already bounded — the surviving candidates are ranked by a *quality* score that averages the two signals most indicative of sequence quality:

$$\text{quality} = \frac{1}{2} \left(\text{PPI} + \frac{\text{No-fold}}{100} \right).$$

The top-ranked candidates of each gene are exported for 3D validation, and the same ranking is used to export larger curated sets (for example the top 25 per gene) together with their sequence-level features.

3.10 Three-dimensional validation

The selected probes are written one per file in the native Boltz-2 input format and submitted to Boltz-2 (Wohlwend et al., 2024), run on a GPU, which predicts the three-dimensional structure of each single-stranded sequence together with confidence estimates: an overall confidence, the **predicted Local Distance Difference Test (pLDDT)** and the **Predicted Aligned Error (PAE)** matrix, in the style of AlphaFold 3 (Abramson et al., 2024). As these confidence metrics are calibrated primarily for proteins,

the overall confidence and the **pLDDT** are used here as relative, independent indicators of the structural soundness of the short probes. The predictions are merged with the sequence metadata to produce a final ranked shortlist combining sequence quality (conservation and No-fold) with the 3D confidence.

3.11 Assembly of probe–graphene models

To obtain atomistic models of the sensing interface, the predicted probe structures were combined with a graphene sheet.

3.11.1 Generation of the graphene sheet with GOPY

The graphene component was built with **GOPY** (Muraru et al., 2020), a free and open-source Python tool for the automated construction of two-dimensional graphene-based molecular models. Building such models by hand is slow and introduces bias, because the position of every atom — and, in functionalised variants, of every chemical group — has to be placed manually; GOPY generates them programmatically instead, which makes larger and more complex sheets practical and reproducible. Besides pristine graphene, the tool can produce graphene oxide and its derivatives (reduced graphene oxide, carboxyl-, epoxy- and hydroxyl-functionalised forms, PEG-aminated reduced graphene oxide and nitrogen-doped graphene), and can introduce vacancy defects (“holes”) in a layer. Each model is written directly in PDB format, which makes it usable by standard structural biology and molecular modelling software without conversion.

A note on how the tool was obtained is warranted, since it conditions reproducibility. GOPY is not published on a package index: it is released by its authors as a standalone Python script under the GNU General Public License, hosted in a public repository,¹ and contains no packaging metadata that would allow installation through `pip`. Retrieving the source directly from that repository and executing it locally is therefore the only available distribution route, and it is the one used here. Care is needed because the name *gopy* is shared by at least two unrelated projects — a Go-to-Python bindings generator and a data-structures package published on PyPI — neither of which has any connection to graphene modelling; the tool used in this work is specifically the one accompanying Muraru et al. (2020), developed by S. Muraru, J. S. Burns and M. Ioniță, to whom credit for it belongs. Only Python 3 with `numpy` and `scipy` is required to run it.

For this work only the pristine graphene functionality was used, since the sensing channel modelled

¹ <https://github.com/Iourarum/GOPY>

here is unfunctionalised graphene. A single layer was generated with the command

```
python GOPY.py generate_PG 120 120 graphene_PG_12nm.pdb
```

where the two numerical arguments are the requested lateral dimensions in Ångström. The resulting sheet was inspected to confirm that it reproduces the expected graphene lattice: it contains 5700 carbon atoms spanning 121.6×119.8 Å, it is strictly planar (all atoms at $z = 0$, with zero dispersion), and the nearest-neighbour carbon–carbon distance is 1.418 Å, the accepted bond length of the graphene honeycomb lattice. Of the 5700 atoms, 5488 are three-coordinated, as expected for the sp^2 interior of the lattice, while the remaining 212 are the two- and one-coordinated atoms that necessarily terminate the edges of a finite, non-periodic sheet. These edge atoms are left unsaturated, as the model is used as a rigid surface for placement rather than in a reactive or dynamical simulation.

The lateral dimensions were not chosen arbitrarily but from the measured footprints of the probes themselves (Section 4.8), so that every probe of the panel fits within the sheet with a margin.

3.11.2 Assembly of the probe–graphene models

Each probe was placed above the sheet by a rigid-body translation, computed with a script built on Biopython (Cock et al., 2009). The probe structure (Boltz-2 CIF) and the sheet (GOPY PDB) are read, and the probe is translated by

$$\Delta = \left(\bar{x}_g - \bar{x}_p, \bar{y}_g - \bar{y}_p, (\bar{z}_g + d) - \min z_p \right),$$

where $\bar{x}, \bar{y}, \bar{z}$ denote centroid coordinates of the sheet (g) and of the probe (p), and d is the chosen separation. The first two components centre the probe over the sheet in the plane; the third places its lowest atom at a height d above the graphene plane.

The separation d was taken from the literature rather than chosen arbitrarily. Van der Waals-corrected density-functional calculations place physisorbed DNA nucleobases at 3.16–3.22 Å above the graphene plane, a regime of strong physisorption dominated by van der Waals interactions (Tayo et al., 2023), consistent with the physisorption picture established for the four bases on graphene and h-BN (Lee et al., 2013). Accordingly, $d = 3.2$ Å was used as the default, so that the assembled model starts at the reported equilibrium contact distance. The separation is an explicit parameter of the script and can be set to any other value.

The probe’s internal conformation and its orientation are left untouched: the operation is a pure translation, so no coordinate is distorted.

The two structures are then merged into a single model and written as one PDB file. Because the PDB format identifies each molecular entity by a single-character chain identifier, and none is reserved for non-biological components such as a graphene layer, an explicit convention was adopted here: the graphene sheet is assigned chain G and the probe chain D, the initials of *graphene* and *DNA*. The choice is arbitrary in the sense that any identifiers would serve, but keeping the two components in separate, consistently named chains is what allows either of them to be selected unambiguously in downstream tools — for instance to hold the sheet rigid during a docking or minimisation run, or to compute probe–surface distances — and the same convention is used in every model produced.

Before writing, the script verifies that no probe atom falls outside the sheet footprint in x – y and refuses to produce a model with overhanging atoms, so that an undersized sheet cannot pass unnoticed. It also reports the shortest probe–sheet interatomic distance, as a check that the placement introduces no steric clash.

It should be stressed that this procedure yields a *starting geometry* — a physically reasonable initial placement of the two components — and not a docked, energy-minimised or dynamically equilibrated structure. It provides the assembled input required by subsequent molecular docking or force-field refinement, but no interaction energy is computed at this stage.

3.12 Integration of a reference probe panel

An existing experimental probe panel was integrated into the same workflow for comparison. Each reference probe, supplied as a (job_name, sequence, species) record, is cleaned (uracil converted to thymine, non-ACGT characters removed) and scored with exactly the same metrics as the designed probes, but by the thresholds of its *own* species and type. Conservation is left empty for these probes, because it is only definable for probes derived from an alignment. The reference probes are stored alongside the designed ones, with a source label, so that the two sets appear side by side with identical columns.

3.13 Diversity and rarefaction analysis

To characterise each target and to justify the sampling depth, an alignment-free analysis based on 4-mer composition is performed on the retrieved sequences. Each sequence is represented by its vector of 4-mer frequencies, and pairwise cosine distances (on the frequency vectors) and Jaccard distances (on 4-mer presence/absence) are computed, together with the fraction of unique sequences, which detects exact duplicates. A rarefaction analysis subsamples N sequences per gene and measures the mean 4-mer

diversity and the richness (number of distinct 4-mers), to determine the value of N beyond which little new information is gained.

3.14 Implementation and reproducibility

All intermediate and final results are written as CSV files. Because **NCBI** does not always return an identical set of records, per-gene counts may vary by a few percent between runs; for strict reproducibility the accession identifiers can be frozen.

Chapter 4

Results and Discussion

This chapter reports the results obtained by applying the pipeline to the six-gene panel, with 100 sequences requested per gene (and a per-gene override of 200 for the data-limited *algD*). It follows the flow of the method: sequence retrieval and curation, the screening funnel, three-dimensional validation, the diversity and rarefaction analyses, the effect of the organism-aware parameters, the comparison with the reference panel, the assembly of probe–graphene models, and the curated dataset of prioritised candidates.

4.1 Retrieval and length curation

Table 4 shows, for each gene, the number of sequences returned by **NCBI** and the number retained after length clustering. The clustering step removed partial and oversized records that would otherwise degrade the alignment; its effect is largest for the genes with the most heterogeneous records. Two genes, *oprL* and *algD*, are intrinsically data-limited in **NCBI** (only 24 and 12 sequences survive respectively), which, as shown below, is reflected in their behaviour through the pipeline.

Gene	Retrieved	After clustering	Alignment (cols)
<i>nuc</i>	99	77	517
<i>rmpM</i>	100	76	1311
<i>lytA</i>	95	87	2594
<i>oprL</i>	27	24	770
<i>algD</i>	19	12	1020
<i>frdB</i>	82	82	489

Table 4: Sequences retrieved and retained per gene, and the resulting alignment width.

4.2 Screening funnel

Across the six genes, the pipeline generated 8455 candidate windows and filtered them down to 2943 that passed both the thermodynamic screen and the secondary-structure screen (Table 5). The melting-temperature and GC filters remove the largest fractions, as expected, since they encode the core hybridisation requirements; the homodimer filter is particularly severe for the GC-rich *P. aeruginosa* genes (*oprL*, *algD*), whose sequences are more prone to self-complementarity. The secondary-structure screen removes only a small additional fraction, confirming that most structurally problematic candidates are already eliminated by the thermodynamic filters.

Gene	Windows	T_m	GC	Hairpin	Basic	Seqfold
<i>nuc</i>	667	326	154	143	143	138
<i>rmpM</i>	1867	1621	1366	1251	841	823
<i>lytA</i>	1833	1274	966	942	820	771
<i>oprL</i>	1360	1273	990	800	349	309
<i>algD</i>	1012	942	763	662	403	389
<i>frdB</i>	1716	932	752	733	522	513
Total	8455	6368	4991	4531	3078	2943

Table 5: Screening funnel: candidate windows surviving each successive filter. “Basic” is the number passing all four thermodynamic criteria; “Seqfold” additionally passes the secondary-structure screen.

4.3 Three-dimensional validation

From the candidates that passed screening, the five highest-quality per gene (30 probes in total) were selected by the quality score and validated with Boltz-2. Of these, 4 were classified as good, 11 as moderate and 15 as low by the overall confidence. The best candidates, strong *both* in sequence quality and in 3D, are listed in Table 6; they come from *algD* and *frdB*. The selected probes are strongly conserved (PPI up to 1.00) and essentially free of self-structure (No-fold above 96), which coincides with the highest structure-prediction confidence and pLDDT among the 30 validated probes. The agreement between the sequence-level signals and the independent 3D prediction provides an internal consistency check on the ranking.

Probe	Gene	PPI	No-fold	Confidence	pLDDT
p5972	<i>algD</i>	0.93	98.4	0.731	0.901
p7716	<i>frdB</i>	0.99	99.6	0.723	0.822
p5938	<i>algD</i>	0.90	96.2	0.721	0.887
p8448	<i>frdB</i>	1.00	98.2	0.716	0.845

Table 6: Top probe candidates, strong in both sequence quality and 3D confidence.

4.4 Genetic diversity of the targets

The alignment-free diversity analysis (Table 7) explains much of the behaviour observed above. *lytA* is by far the most diverse gene (mean cosine distance 0.187 and 93% unique sequences), consistent with its known allelic diversity in *S. pneumoniae* (Obregón et al., 2002); this justifies the relaxed conservation threshold applied to it. At the opposite extreme, *frdB* is the most conserved and most redundant gene (mean cosine distance 0.025 and only 35% unique sequences, i.e. many exact duplicates in NCBI), which is why it yields probes of very high conservation. *algD* has 100% unique sequences but very few of them, which accounts for its instability at low sampling and motivates the automatic length-clustering and the per-gene allowance of more sequences. The cosine and Jaccard distances are complementary — the former weighs 4-mer frequencies, the latter only their presence or absence — and they concur in identifying *frdB* as the most conserved and redundant gene, giving an alignment-independent reading of diversity.

Gene	<i>n</i>	Cosine (mean)	Cosine (max)	Jaccard (mean)	Unique (%)
<i>nuc</i>	77	0.059	0.425	0.163	61.0
<i>rmpM</i>	76	0.079	0.247	0.069	50.0
<i>lytA</i>	87	0.187	0.440	0.125	93.1
<i>oprL</i>	24	0.050	0.358	0.096	91.7
<i>algD</i>	12	0.097	0.332	0.117	100.0
<i>frdB</i>	82	0.025	0.052	0.061	35.4

Table 7: Alignment-free 4-mer diversity per gene. Bold marks the most diverse (*lytA*) and the most conserved/redundant (*frdB*) genes.




Figure 2: Pairwise 4-mer distance heatmaps per gene (alignment-free). Each cell is the 4-mer distance between two retrieved sequences of the gene.

4.5 Choice of sampling depth

The rarefaction analysis showed that both the 4-mer diversity and the 4-mer richness saturate at around $N \approx 50$ sequences for the most demanding gene (*nuc*), and much earlier for the others. Requesting up to 100 sequences per gene therefore captures essentially all of the available diversity with a comfortable margin, while additional sequences add little new information. The true limiting factor is not this cap but the small number of records available in [NCBI](#) for the data-limited genes.




Figure 3: Rarefaction of the mean 4-mer diversity (left) and of the richness, i.e. the number of distinct 4-mers (right), per gene, as a function of the number of subsampled sequences.

4.6 Effect of the organism-aware parameters

The value of resolving thresholds per organism type and species is clearest when scoring the reference panel, which spans bacteria, viruses, a fungus, a protozoan and a human control. When the 74 reference probes were scored with a single global (bacterial) criterion, 29 passed the basic screen; when instead each

was scored by the criteria of its own species and type, 39 passed. The additional probes are precisely those whose genome composition differs from the balanced bacterial case — for example the AT-rich *Plasmodium* probe, which the global GC window would reject but which is appropriate under the protozoan profile. This confirms that organism-aware thresholds avoid unfair rejections without loosening the criteria for the balanced targets.

4.7 Comparison with the reference panel

Scored with identical metrics, the designed probes and the reference panel can be compared directly. Among the probes that pass the basic screen, the designed probes have a higher mean melting temperature (59.3 versus 56.9 °C) and a very similar mean GC content (0.52 versus 0.53), while the reference probes have a marginally higher mean No-fold score (86.0 versus 82.9). Because both sets carry a source label and share every column, they can also be submitted together to Boltz-2, enabling a like-for-like comparison in three dimensions as well as in sequence space.

4.8 Probe–graphene models

All 30 probes validated in 3D were combined with a graphene sheet to obtain atomistic models of the sensing interface (Section 3.11). The predicted probe structures contain between 359 and 585 atoms (median 416), consistent with single strands of 18–28 nucleotides.

Sizing the sheet. The lateral dimensions of the sheet were set from the probes themselves. Measuring, for each predicted structure, its largest extent in the plane gives a median of 42.0 Å but a range that reaches 82.9 Å (Table 8): the probes are predicted in a variety of extended and compact conformations, and the most extended are nearly twice the median. A 4 × 4 nm sheet would therefore have been too small for 16 of the 30 probes, whose atoms would overhang the edges — an artefact that the fit check built into the assembly script rejects. A 121.6 × 119.8 Å sheet (5700 carbon atoms) was used instead, which accommodates all 30 probes, the most extended one still leaving a margin of 19.3 Å on each side.

Assembled models. The complete set of 30 probe–graphene models was built, each comprising 6059–6285 atoms. Table 8 summarises their geometry by gene. Two observations follow. First, no model contains a steric clash: the closest probe–sheet atom pair lies between 3.20 and 3.47 Å (median 3.25 Å) across all 30 systems, that is, within the physisorption regime reported for nucleobases on

graphene (Tayo et al., 2023) and never below it. The small spread above the nominal 3.2 Å separation arises because the probe atom closest to the surface is generally offset laterally from the graphene atoms beneath it, so the true interatomic distance slightly exceeds the vertical one.

Second, the probes stand well above the surface: their vertical extent ranges from 20.6 to 80.4 Å. This is relevant to the transduction context of Chapter 2, since in an electrolyte-gated device the charge carried far from the graphene plane is screened, and only the portion of the probe within the Debye length contributes appreciably to the signal (Purwidyantri et al., 2021). The models therefore make explicit a property that the sequence-level metrics cannot capture: probes of identical length may present very different amounts of their structure close to the sensing surface.

Gene	Models	Median extent (Å)	Max. extent (Å)	Height (Å)	Closest contact (Å)
<i>nuc</i>	5	45.0	82.9	23.0–80.4	3.21–3.40
<i>rmpM</i>	5	54.9	56.1	20.6–48.2	3.22–3.27
<i>lytA</i>	5	38.2	66.7	23.0–47.5	3.24–3.35
<i>oprL</i>	5	46.6	63.2	25.4–44.2	3.20–3.45
<i>algD</i>	5	32.9	47.7	20.6–42.9	3.26–3.38
<i>frdB</i>	5	39.0	55.6	27.7–51.4	3.20–3.47
All	30	42.0	82.9	20.6–80.4	3.20–3.47

Table 8: Geometry of the 30 assembled probe–graphene models. “Extent” is the largest in-plane dimension of the probe (which determined the sheet size), “height” its vertical extent above the sheet, and “closest contact” the shortest probe–sheet interatomic distance.

Figure 4 shows the assembled model of the highest-confidence probe of each of the two best-performing genes, drawn directly from the atomic coordinates of the output files. The side view makes the assembly geometry explicit — the planar graphene lattice, the probe resting above it, and the 3.2 Å separation — while the top view shows the probe footprint contained well within the sheet. The two panels also illustrate the conformational contrast noted above: *p5972* is compact, reaching about 20 Å in height, whereas *p7716* is extended and reaches beyond 40 Å, so a much larger fraction of its structure lies far from the surface.




Figure 4: Assembled probe-graphene models for *p5972* (*algD*, top) and *p7716* (*frdB*, bottom), rendered from the model coordinates. Left: side view, with the graphene plane, the probe above it and the 3.2 Å separation indicated. Right: top view, showing the probe footprint within the 121.6×119.8 Å sheet. Probe atoms are coloured by element; graphene carbons are grey.

These assembled structures are starting geometries for subsequent docking or force-field refinement. No interaction energy is computed at this stage, so they support no claim about binding affinity or preferred orientation; what they establish is that the whole validated panel can be assembled into clash-free, physically sensible interface models by an automated procedure.

What the models already indicate

Although the models are unrefined, they make one device-relevant quantity measurable that no sequence-level metric captures: how much of the probe sits close enough to the graphene to be electrically visible. In

an electrolyte-gated device the signal from charges beyond the Debye length is screened, and that length depends strongly on the buffer: for the electrolytes characterised on graphene DNA transistors it ranges from 0.76 nm in $1\times$ PBS to 2.41 nm in $0.1\times$ PBS and 7.61 nm in $0.01\times$ PBS (Purwidyantri et al., 2021). Taking the first two as reference, the fraction of probe atoms lying within the Debye length of the sheet varies markedly across the panel: from 2.0 to 14.3% (median 6.9%) at $1\times$ PBS, and from 22 to 100% at $0.1\times$ PBS. In other words, at physiological ionic strength the great majority of every probe would be screened, whereas in dilute buffer some probes become almost fully visible and others remain largely hidden — a spread that is a property of their predicted conformation, not of their sequence.

Two consequences follow for the candidates themselves. First, this geometric criterion is *not* aligned with the 3D confidence used to rank them: across the 30 models the correlation between the Boltz-2 confidence and the fraction of the probe within 0.76 nm is weak and slightly negative ($r = -0.19$). A probe can therefore be well predicted and still present little of itself to the sensing surface. Second, and more usefully, the two criteria can be combined. The probe *p5972 (algD)* is the only candidate that is simultaneously the best-scoring in 3D (confidence 0.731, good) and among the most compact, with a height of 20.6 Å and *all* of its atoms within 2.41 nm of the surface; *p2392 (rmpM)* shows the same favourable geometry with a more modest confidence. By contrast *p7716* and *p8448 (frdB)*, despite ranking good in 3D, are extended structures with only 38 and 22% of their atoms within 2.41 nm, so a large part of them would fall outside the sensing volume.

This analysis carries an important caveat. The probe orientation in these models is the one predicted by Boltz-2, preserved by the rigid-body placement and *not* optimised against the surface; since single-stranded DNA is known to adsorb flat on graphene through π - π stacking (Green and Norton, 2015), the orientations adopted after docking or dynamics will differ, and with them the exact coverage values. The numbers above should therefore be read as an indication of the range of behaviour the panel can display, and as evidence that interface geometry is an independent axis of probe quality worth carrying into the selection, rather than as predictions of device response.

4.9 Curated dataset of prioritised candidates

Finally, the same quality ranking was used to export a larger curated dataset: the top 25 probes per gene (150 probes in total), each with its sequence-level features (T_m , GC, conservation, No-fold and folding energies). This dataset extends the set of 30 probes validated in 3D to a broader, biophysically annotated selection of the highest-quality candidates per gene.

Chapter 5

Conclusions and Future Work

5.1 Conclusions

This dissertation set out to design, *in silico*, high-quality **ssDNA** capture probes for **GFET** biosensors targeting a panel of clinically relevant bacterial pathogens. The main outcome is a single, reproducible computational pipeline that turns a target gene into a small set of ranked, biophysically validated probe candidates, with every design decision grounded in the literature.

The pipeline integrates, in one automated workflow, the stages that are usually performed separately or by hand: retrieval and curation of sequences from **NCBI**, automatic length clustering, multiple sequence alignment and **PPI**-based conservation scoring, nearest-neighbour thermodynamic screening, secondary-structure evaluation with a continuous No-fold score, and a final three-dimensional validation with Boltz-2. Two design choices give the method its generality and rigour. First, thresholds are resolved through a layered, *organism-aware* scheme (gene → species → type → global), so that pathogens with very different genome composition are judged by fair criteria; this was shown to recover candidates — such as AT-rich probes — that a single global rule would wrongly discard. Second, candidates are prioritised by a transparent quality score combining conservation and the No-fold score, which concentrates the expensive 3D validation on the most promising sequences.

Applied to the six-gene panel, the pipeline reduced 8455 candidate windows to 2943 that passed screening and identified a shortlist of top candidates whose high sequence quality was independently corroborated by the highest 3D confidence and **pLDDT**. Alignment-free diversity and rarefaction analyses characterised each target and justified the sampling depth in a principled, non-arbitrary way. An existing experimental probe panel was folded into the same workflow and scored with identical metrics, enabling a direct comparison, and a curated 150-probe dataset of the highest-quality candidates per gene was produced.

The validated probes were further assembled with a graphene sheet into atomistic models of the sens-

ing interface, using a pristine layer built with GOPY and a placement set at the equilibrium physisorption distance reported for nucleobases on graphene. All 30 models are free of steric clashes. Beyond demonstrating that the panel can be assembled automatically, these models exposed a selection criterion that the sequence-level analysis cannot reach: the fraction of each probe lying within the Debye length of the surface varies from 2 to 14% at physiological ionic strength and from 22 to 100% in dilute buffer, and it is essentially uncorrelated with the 3D confidence. Interface geometry is thus an independent axis of probe quality. On the two criteria combined, *p5972 (algD)* stands out, being both the highest-confidence structure and one of the most compact.

Overall, the work delivers a solid, extensible foundation for the computational stage of **GFET** biosensor development: it produces prioritised, structurally sound probe candidates and generalises to new genes and organisms with minimal configuration.

5.2 Prospects for future work

Several directions follow naturally from this work:

- **Experimental validation.** The ultimate test of the designed probes is their performance on fabricated **GFET** devices; the ranked shortlist provides the immediate candidates for functionalisation and characterisation.
- **Specificity screening.** Adding an automated specificity check (for example a BLAST step against non-target genomes and the host) would complement the current conservation-driven, sensitivity-oriented selection.
- **Comparison of 3D methods.** The Boltz-2 predictions can be benchmarked against an independent molecular-docking pipeline, to assess the agreement of the two structural views of the probe–target interaction.
- **Data-driven modelling.** The curated feature dataset opens the way to a machine-learning model that predicts probe quality directly, potentially accelerating or refining the ranking.
- **Broader panels.** As the framework already supports viral, fungal, protozoan and host profiles, extending the panel to further genes and non-bacterial pathogens is largely a matter of configuration.

Bibliography

Josh Abramson, Jonas Adler, Jack Dunger, et al. Accurate structure prediction of biomolecular interactions with AlphaFold 3. *Nature*, 630(8016):493–500, 2024. doi: 10.1038/s41586-024-07487-w.

O. G. Brakstad, K. Aasbakk, and J. A. Maeland. Detection of *staphylococcus aureus* by polymerase chain reaction amplification of the *nuc* gene. *Journal of Clinical Microbiology*, 30(7):1654–1660, 1992. doi: 10.1128/jcm.30.7.1654-1660.1992.

Rui Campos, Jérôme Borne, Joana Rafaela Guerreiro, George Machado Jr., Maria Fátima Cerqueira, Dmitri Y. Petrovykh, and Pedro Alpuim. Attomolar label-free detection of DNA hybridization with electrolyte-gated graphene field-effect transistors. *ACS Sensors*, 4(2):286–293, 2019. doi: 10.1021/acssensors.8b00344.

Maria da Glória S. Carvalho, Maria L. Tondella, Karen McCaustland, et al. Evaluation and improvement of real-time PCR assays targeting *lyta*, *ply*, and *psaa* genes for detection of pneumococcal DNA. *Journal of Clinical Microbiology*, 45(8):2460–2466, 2007. doi: 10.1128/JCM.02498-06.

Peter J. A. Cock, Tiago Antao, Jeffrey T. Chang, Brad A. Chapman, Cymon J. Cox, Andrew Dalke, Iddo Friedberg, Thomas Hamelryck, Frank Kauff, Bartek Wilczynski, and Michiel J. L. de Hoon. Biopython: Freely available Python tools for computational molecular biology and bioinformatics. *Bioinformatics*, 25(11):1422–1423, 2009. doi: 10.1093/bioinformatics/btp163.

D. De Vos, A. Lim Jr., J.-P. Pirnay, M. Struelens, C. Vandenvelde, L. Duinslaeger, A. Vanderkelen, and P. Cornelis. Direct detection and identification of *pseudomonas aeruginosa* in clinical samples such as skin biopsy specimens and expectorations by multiplex PCR based on two outer membrane lipoprotein genes, *opri* and *oprl*. *Journal of Clinical Microbiology*, 35(6):1295–1299, 1997. doi: 10.1128/jcm.35.6.1295-1299.1997.

Andrew D. Ellington and Jack W. Szostak. In vitro selection of RNA molecules that bind specific ligands. *Nature*, 346(6287):818–822, 1990. doi: 10.1038/346818a0.

- Malcolm J. Gardner, Neil Hall, Eula Fung, et al. Genome sequence of the human malaria parasite *Plasmodium falciparum*. *Nature*, 419(6906):498–511, 2002. doi: 10.1038/nature01097.
- N. S. Green and M. L. Norton. Interactions of DNA with graphene and sensing applications of graphene field-effect transistor devices: A review. *Analytica Chimica Acta*, 853:127–142, 2015. doi: 10.1016/j.aca.2014.10.038.
- Changchun Hu, Shuting Yang, Shuo Li, Xueying Liu, Yuan Liu, Zhu Chen, Hong Chen, Ning He, Hong Cui, and Yulin Deng. Viral aptamer screening and aptamer-based biosensors for virus detection: A review. *International Journal of Biological Macromolecules*, 276:133935, 2024. doi: 10.1016/j.ijbiomac.2024.133935.
- Michael T. Hwang, Mohammad Heiranian, Yerim Kim, Seungyong You, Juyoung Leem, Amir Taqieddin, Vahid Faramarzi, Yanfeng Jing, Insu Park, Arend M. van der Zande, Sungwoo Nam, Narayana R. Aluru, and Rashid Bashir. Ultrasensitive detection of nucleic acids using deformed graphene channel field effect biosensors. *Nature Communications*, 11:1543, 2020. doi: 10.1038/s41467-020-15330-9.
- Integrated DNA Technologies. OligoAnalyzer tool. <https://www.idtdna.com/pages/tools/oligoanalyzer>. Online oligonucleotide analysis tool, accessed 2026.
- John Jumper, Richard Evans, Alexander Pritzel, et al. Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873):583–589, 2021. doi: 10.1038/s41586-021-03819-2.
- Kazutaka Katoh and Daron M. Standley. MAFFT multiple sequence alignment software version 7: Improvements in performance and usability. *Molecular Biology and Evolution*, 30(4):772–780, 2013. doi: 10.1093/molbev/mst010.
- C. Kolm, I. Cervenka, C. Aschl, N. Baumann, S. Jakwerth, R. Krska, R. L. Mach, R. Sommer, M. C. DeRosa, A. K. T. Kirschner, A. H. Farnleitner, and G. H. Reischer. DNA aptamers against bacterial cells can be efficiently selected by a SELEX process using state-of-the-art qPCR and ultra-deep sequencing. *Scientific Reports*, 10:20917, 2020. doi: 10.1038/s41598-020-77221-9.
- Thomas R. Kozel and Amanda R. Burnham-Marusich. Point-of-care testing for infectious diseases: Past, present, and future. *Journal of Clinical Microbiology*, 55(8):2313–2320, 2017. doi: 10.1128/JCM.00476-17.

- Siva Kumar Krishnan, Nandini Nataraj, M. Meyyappan, and Umapada Pal. Graphene-based field-effect transistors in biosensing and neural interfacing applications: Recent advances and prospects. *Analytical Chemistry*, 95(5):2590–2622, 2023. doi: 10.1021/acs.analchem.2c03399.
- Jun-Ho Lee, Yun-Ki Choi, Hyun-Jung Kim, Ralph H. Scheicher, and Jun-Hyung Cho. Physisorption of DNA nucleobases on h-BN and graphene: vdW-corrected DFT calculations. *The Journal of Physical Chemistry C*, 117(26):13435–13441, 2013. doi: 10.1021/jp402403f.
- Ana Lima. ViruScope: Conservation (PPI) and no-fold scoring for nucleic acid probe design. <https://github.com/anasfplima/ViruScope>, 2023. Software repository, accessed 2026.
- Chun-Hua Lu, Huang-Hao Yang, Chun-Ling Zhu, Xi Chen, and Guo-Nan Chen. A graphene platform for sensing biomolecules. *Angewandte Chemie International Edition*, 48(26):4785–4787, 2009. doi: 10.1002/anie.200901479.
- Sebastian Muraru, Jorge S. Burns, and Mariana Ioniță. GOPY: A tool for building 2D graphene-based computational models. *SoftwareX*, 12:100586, 2020. doi: 10.1016/j.softx.2020.100586.
- Christopher J. L. Murray, Kevin S. Ikuta, Fablina Sharara, et al. Global burden of bacterial antimicrobial resistance in 2019: A systematic analysis. *The Lancet*, 399(10325):629–655, 2022. doi: 10.1016/S0140-6736(21)02724-0.
- Nako Nakatsuka, Kyung-Ae Yang, John M. Abendroth, et al. Aptamer–field-effect transistors overcome debye length limitations for small-molecule sensing. *Science*, 362(6412):319–324, 2018. doi: 10.1126/science.aao6750.
- William C. Nierman, Arnab Pain, Michael J. Anderson, et al. Genomic sequence of the pathogenic and allergenic filamentous fungus *aspergillus fumigatus* Af293. *Nature*, 438(7071):1151–1156, 2005. doi: 10.1038/nature04332.
- K. S. Novoselov, A. K. Geim, S. V. Morozov, D. Jiang, Y. Zhang, S. V. Dubonos, I. V. Grigorieva, and A. A. Firsov. Electric field effect in atomically thin carbon films. *Science*, 306(5696):666–669, 2004. doi: 10.1126/science.1102896.
- V. Obregón, J. L. García, E. García, A. Fenoll, R. López, and P. García. Molecular peculiarities of the *lytA* gene isolated from clinical pneumococcal strains that are bile insoluble. *Journal of Clinical Microbiology*, 40(7):2545–2554, 2002. doi: 10.1128/JCM.40.7.2545-2554.2002.

- Jinglei Ping, Ramya Vishnubhotla, Amey Vrudhula, and A. T. Charlie Johnson. Scalable production of high-sensitivity, label-free DNA biosensors based on back-gated graphene field-effect transistors. *ACS Nano*, 10(9):8700–8704, 2016. doi: 10.1021/acsnano.6b04110.
- Agustin Purwidyantri, Tiago Domingues, Jérôme Borme, Joana Rafaela Guerreiro, Andrey Ipatov, Cláudia M. Abreu, Manuel Martins, Pedro Alpuim, and Marta Prado. Influence of the electrolyte salt concentration on DNA detection with graphene transistors. *Biosensors*, 11(1):24, 2021. doi: 10.3390/bios11010024.
- Agustin Purwidyantri, Andrey Ipatov, Tiago Domingues, Jérôme Borme, Manuel Martins, Pedro Alpuim, and Marta Prado. Programmable graphene-based microfluidic sensor for DNA detection. *Sensors and Actuators B: Chemical*, 367:132044, 2022. doi: 10.1016/j.snb.2022.132044.
- John SantaLucia Jr. and Donald Hicks. The thermodynamics of DNA structural motifs. *Annual Review of Biophysics and Biomolecular Structure*, 33:415–440, 2004. doi: 10.1146/annurev.biophys.32.110601.141800.
- F. Schedin, A. K. Geim, S. V. Morozov, E. W. Hill, P. Blake, M. I. Katsnelson, and K. S. Novoselov. Detection of individual gas molecules adsorbed on graphene. *Nature Materials*, 6(9):652–655, 2007. doi: 10.1038/nmat1967.
- C. Stephens, N. M. Goodey, and U. Gubler. A beginner’s guide to SELEX and DNA aptamers. *Analytical Biochemistry*, 703:115890, 2025. doi: 10.1016/j.ab.2025.115890.
- C. K. Stover, X. Q. Pham, A. L. Erwin, S. D. Mizoguchi, et al. Complete genome sequence of *pseudomonas aeruginosa* PAO1, an opportunistic pathogen. *Nature*, 406(6799):959–964, 2000. doi: 10.1038/35023079.
- Di Sun, Miao Sun, Jialu Zhang, Xin-Hui Lin, Yinkun Zhang, Fanghe Lin, Peng Zhang, Chaoyong James Yang, and Jia Song. Computational tools for aptamer identification and optimization. *TrAC Trends in Analytical Chemistry*, 157:116767, 2022. doi: 10.1016/j.trac.2022.116767.
- Benjamin O. Tayo, Michael A. Walkup, and Serkan Caliskan. Adsorption of DNA nucleobases on single-layer Ti_3C_2 MXene and graphene: vdW-corrected DFT and NEGF studies. *AIP Advances*, 13(8):085213, 2023. doi: 10.1063/5.0159563.
- Joshua Timmons. seqfold: Nucleic acid secondary structure prediction in Python. <https://github.com/Lattice-Automation/seqfold>, 2020. Software; nearest-neighbour free-energy model.

Craig Tuerk and Larry Gold. Systematic evolution of ligands by exponential enrichment: RNA ligands to bacteriophage T4 DNA polymerase. *Science*, 249(4968):505–510, 1990. doi: 10.1126/science.2200121.

Douglas H. Turner and David H. Mathews. NNDB: The nearest neighbor parameter database for predicting stability of nucleic acid secondary structure. *Nucleic Acids Research*, 38(suppl 1):D280–D282, 2010. doi: 10.1093/nar/gkp892.

Andreas Untergasser, Ioana Cutcutache, Triinu Koressaar, Jian Ye, Brant C. Faircloth, Maido Remm, and Steven G. Rozen. Primer3—new capabilities and interfaces. *Nucleic Acids Research*, 40(15):e115, 2012. doi: 10.1093/nar/gks596.

James G. Wetmur. DNA probes: Applications of the principles of nucleic acid hybridization. *Critical Reviews in Biochemistry and Molecular Biology*, 26(3-4):227–259, 1991. doi: 10.3109/10409239109114069.

Jeremy Wohlwend, Gabriele Corso, Saro Passaro, et al. Boltz-1: Democratizing biomolecular interaction modeling. *bioRxiv*, 2024. doi: 10.1101/2024.11.19.624167. Preprint.

World Health Organization. WHO bacterial priority pathogens list, 2024. Technical report, World Health Organization, Geneva, 2024.

Joseph N. Zadeh, Conrad D. Steenberg, Justin S. Bois, Brian R. Wolfe, Marshall B. Pierce, Asif R. Khan, Robert M. Dirks, and Niles A. Pierce. NUPACK: Analysis and design of nucleic acid systems. *Journal of Computational Chemistry*, 32(1):170–173, 2011. doi: 10.1002/jcc.21596.

Yingyue Zhu, Yilin Cai, Liguang Xu, Lixue Zheng, Limei Wang, Bin Qi, and Chuanlai Xu. Building an aptamer/graphene oxide FRET biosensor for one-step detection of bisphenol A. *ACS Applied Materials & Interfaces*, 7(14):7492–7496, 2015. doi: 10.1021/acsami.5b00199.

Place here information about funding, FCT project, etc. in which the work is framed. Leave empty otherwise.